

# PS2 - Fashion MNIST

October 9, 2018

## 1 Classify the kind of clothing item (pant, shirt, etc.) shown in the Fashion MNIST dataset

- Do this using the best practices discussed in class (i.e. one hot encoding, ReLU, CNNs when you should, etc.)
- Have reasonable hyperparameters

```
In [16]: # Importing the libraries
import numpy as np
import keras
from matplotlib import pyplot as plt
import seaborn as sns
from keras.models import Sequential
from keras.layers import Dense, Dropout, Activation, Flatten, BatchNormalization
from keras.layers import Conv2D, MaxPooling2D
from sklearn.metrics import confusion_matrix
from sklearn.model_selection import train_test_split
# from keras import backend as K
# K.tensorflow_backend._get_available_gpus()
```

## 2 Loading the Training and Testing Data


```
In [2]: # Loading the Training and Testing Data
trainImages = np.load('./Fashion MNIST/trainImages.npy')
testImages = np.load('./Fashion MNIST/testImages.npy')
trainLabels = np.load('./Fashion MNIST/trainLabels.npy')
testLabels = np.load('./Fashion MNIST/testLabels.npy')
```

## 3 Converting to Float32 to do input data Normalization

```
In [3]: # Converting to Float32 to do input data Normalization
trainImages = trainImages.astype('float32') #images loaded in as int64, 0 to 255 integers
testImages = testImages.astype('float32')
# Normalization
trainImages /= 255
testImages /= 255
```

## 4 Showing the Input Data after Normalizing

```
In [4]: # Preview the training data
plt.figure(figsize=(12,10))# Showing the Input Data after Normalizing
x, y = 3, 3
for i in range(6):
    plt.subplot(y, x, i+1)
    plt.imshow(trainImages[i].reshape((28,28)),interpolation='nearest')
plt.show()
```



## 5 Splitting the Data into Training and Validation Data

```
In [5]: # Splitting the Training Data into training and Validation Data
X_train,X_val, Y_train,Y_val = train_test_split(trainImages, trainLabels, test_size=0.2)

print("For this problem the number of training data is:" + str(X_train.shape[0]) + "\nThe
" validation data is:" + str(X_val.shape[0]) + "\nand the number of test data is:" + str
```

For this problem the number of training data is:48000  
The number of validation data is:12000  
and the number of test data is:10000

## 6 Setting Hyper Paramters

```
In [6]: # Hyper Parameters
        batch_size = 64
        num_classes = testLabels.shape[1]
        epochs = 20
```

## 7 Creating Model

```
In [7]: # create model
        img_rows, img_cols = 28, 28
        input_shape = (1,img_rows, img_cols)
        model = Sequential()
        model.add(Conv2D(32, kernel_size=(3, 3),
                        activation='relu',data_format='channels_first',use_bias=True ,
                        # bias_initializer = keras.initializers.glorot_uniform(seed=None),
                        input_shape=input_shape)) # Adding weights for tuning
        # model.add(BatchNormalization())
        model.add(Conv2D(64, (3, 3), activation='relu')) # Added another layer as this
                                                         # smoothens the training accuracy cu

        # model.add(BatchNormalization())
        model.add(MaxPooling2D(pool_size=(2, 2))) # Reducing the parmeters
        model.add(Dropout(0.25)) # Adding non linearity and randomness
        model.add(Flatten()) # Converting to 1D as required before
                             # fully connected layer.

        model.add(Dense(128, activation='relu')) # fully connected layer with 128 hid
        # model.add(BatchNormalization())
        model.add(Dropout(0.5)) # Adding non linearity and randomness
        model.add(Dense(num_classes, activation='softmax')) # Final fully connected layer.
        # Compile model
        model.compile(loss=keras.losses.categorical_crossentropy,
                      optimizer=keras.optimizers.Adadelta(),metrics=['accuracy'])
        # Get Model Summary
        model.summary()
```

| Layer (type)                 | Output Shape       | Param # |
|------------------------------|--------------------|---------|
| conv2d_1 (Conv2D)            | (None, 32, 26, 26) | 320     |
| conv2d_2 (Conv2D)            | (None, 30, 24, 64) | 15040   |
| max_pooling2d_1 (MaxPooling2 | (None, 15, 12, 64) | 0       |
| dropout_1 (Dropout)          | (None, 15, 12, 64) | 0       |
| flatten_1 (Flatten)          | (None, 11520)      | 0       |

|                             |             |         |
|-----------------------------|-------------|---------|
| dense_1 (Dense)             | (None, 128) | 1474688 |
| -----                       |             |         |
| dropout_2 (Dropout)         | (None, 128) | 0       |
| -----                       |             |         |
| dense_2 (Dense)             | (None, 10)  | 1290    |
| =====                       |             |         |
| Total params: 1,491,338     |             |         |
| Trainable params: 1,491,338 |             |         |
| Non-trainable params: 0     |             |         |
| -----                       |             |         |

## 8 Weights and Biases for First Layer

```
In [8]: # Before Training
weights = model.layers[0].get_weights()
w0 = np.array(weights[0])
b0 = np.array(weights[1])
print("The weights for first layer has dimensions of " + str(w0.shape))
print()
print("The Biases for first layer has dimensions of " + str(b0.shape)+" and the values a")
print()
print("Just so we know that the weights have changed after training,\n
we will be comparing the sum of weights of he first layer before and after training.")
print()
print("The sum of weights of the first layer is :" + str(w0.sum()))
print("The sum of biases of the first layer is :" + str(b0.sum()))
```

The weights for first layer has dimensions of (3, 3, 1, 32)

[illegible]

Just so we know that the weights have changed after training, we will be comparing the sum of we

```
The sum of weights of the first layer is :-0.47791335
The sum of biases of the first layer is :0.0
```

## 9 Training the Model

```
In [9]: # Train
        # Here I have considered the test data as the validation data.
        h = model.fit(X_train, Y_train,
                       batch_size=batch_size,
                       epochs=epochs,
```

```

        verbose=0,
        validation_data=(X_val, Y_val))

# Evaluate Accuracy

score = model.evaluate(testImages, testLabels, verbose=0)
print('Training Loss:', h.history['loss'][-1])
print('Training Accuracy:', h.history['acc'][-1])
print()
print('Validation Loss:', h.history['val_loss'][-1])
print('Validation Accuracy:', h.history['val_acc'][-1])
print()
print('Test Loss:', score[0])
print('Test Accuracy:', score[1])
print()

# Plot Graphs
# print(h.history.keys())
accuracy = h.history['acc']
val_accuracy = h.history['val_acc']
loss = h.history['loss']
val_loss = h.history['val_loss']
epochs = range(len(accuracy))
plt.plot(epochs, accuracy, 'r', label='Training accuracy')
plt.plot(epochs, val_accuracy, 'b', label='Validation accuracy')
plt.title('Test accuracy : ' + str(score[1]))
plt.legend()
plt.show()
plt.figure()
plt.plot(epochs, loss, 'r', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Test loss: '+str(score[0]))
plt.legend()
plt.show()

```

Training Loss: 0.13424885125209887



Training Accuracy: 0.951625

Validation Loss: 0.2452846585114797

Validation Accuracy: 0.923

Test Loss: 0.2743830148816109

Test Accuracy: 0.9164



## 10 Inferences

- The same model was used as used in MNIST. The model seems to be working good with only increase in number of Epochs to 20. the batch size was reduced to 64 for more updates.
- Before using the same model I tried Fashion MNIST with other models as well. Below are the accuracy received with each model:
  - With 3 Dense (fully connected layers); Test Accuracy after 12 epochs was; 90.82 %
  - With 2 Dense and 1 convolution : Test Accuracy after 12 epochs: 91.11 %
  - With 2 Dense and 2 convolution: Test Accuracy is after 12 epochs : 91.74 %
  - Therefore, the final option was adopted.
- Therefore, when there is image data, use of conv2D layers help in feature extraction.
- There are 2 Convolution Layers and 2 Dense or Fully Connected Layers with subsequent Max Pooling and Dropout Layer.
- The number of hidden units in each layer are experimentally fine tuned. The units were experimented in  $2^n$  multiples like 16, 32, 64, 128, 256.
- **\*\* Input Conv2D layer: \*\*** The weights are initialized by Glorot uniform initializer, which is the default. This is set as the values are nicely zero centered and within -1 and 1. The biases are set to zero as it improves the accuracy by 0.14 as compared to randomly distributed values between (-1 and 1) (found out by experimentation). This is also recommended as per in CS231 from Stanford.
  - The kernel size is set to minimum of (3,3). This is because the image size is already small (28,28) and by using large kernel size there won't be enough data left to add more convolution layers.
  - the activation is set to RELU, as it is the most recommended and also as the input is normalized between 0 and 1.
- **\*\* Max Pooling layer: \*\*** for pooling Max Pooling is used. Again as the input size is very less. The maxpooling kernel size is set to (2,2). This works as there is still enough information in a (2,2) window and it also does not throw out a lot of valuable data.
- **\*\* Drop Out: \*\*** This is very Essential as it helps prevent over fitting. Two dropout layers are used.
  - The first one is in between convolution layers. The dropout of 0.25 was determined to be optimal. At this value the convergence of Training accuracy to validation accuracy was smooth and gradual as compared to other values.
  - The second dropout was used just before the final fully connected layer. This is kept at 0.45 as experimentally determined and is used to reduce overfitting.
- **\*\* Dense Layers: \*\*** Two fully connect layers are used with a dropout in between. This seems to give good performances in MNIST and Fashion MNIST.
  - For final fully connected layer, softmax is used as activation. This is because there are a total of 10 classes and the predicted output can exactly be any one of these classes.
- **\*\* Loss \*\*** : as one hot encoding is used, categorical\_crossentropy loss is used as loss function. This is used as it is the most recommended one.
- **\*\* Optimizer \*\*** : Different Optimizers were tried out. Even for Fashion MNIST, for different optimizers there was not a lot of change in the convergence time or the number of epochs.

- For example, I used Adadelata and each epoch took 4.5 seconds less to finish as compared to Adam, also the final test accuracy was higher for Adadelata. Therefore I finally decided to use Adadelata as my optimizers. After playing with the parmaters, I come to the conclusion that the default value works the best.

## 11 Weights and Biases for First Layer After Training

```
In [13]: # After Training
weights = model.layers[0].get_weights()
w0 = np.array(weights[0])
b0 = np.array(weights[1])
print("The weights for first layer has dimensions of " + str(w0.shape))
print()
print("The Biases for first layer has dimensions of " + str(b0.shape)+" and the values")
print()
print("The sum of weights of the first layer is :" + str(w0.sum()))
print("The sum of biases of the first layer is :" + str(b0.sum()))
print("As we can see the sum has changed, indicating that the weights are now tuned.")
```

The weights for first layer has dimensions of (3, 3, 1, 32)

The Biases for first layer has dimensions of (32,) and the values are :

```
[ 0.03732191  0.03564902 -0.02086273  0.00827906 -0.01473734  0.01437863
 -0.00281613  0.00623995 -0.00739944 -0.06542691 -0.00145124 -0.14998387
 -0.01268976 -0.01216698 -0.01292941  0.02989011  0.01361851  0.01142375
  0.01032185 -0.03882701 -0.01248314  0.02152263 -0.09115795 -0.00246029
 -0.00022315 -0.00455045  0.04222413  0.00285289 -0.02125233  0.00818023
  0.05506296 -0.07067282]
```

The sum of weights of the first layer is :-0.62720954

The sum of biases of the first layer is :-0.24512535

As we can see the sum has changed, indicating that the weights are now tuned.

## 12 Calculate Confusion Matrix

```
In [14]: # Predict the values from the validation dataset
Y_pred = model.predict(testImages)
# Convert predictions classes to one hot vectors
Y_pred_classes = np.argmax(Y_pred, axis = 1)
# Convert validation observations to one hot vectors
Y_true = np.argmax(testLabels, axis = 1)
# compute the confusion matrix
confusion_mtx = confusion_matrix(Y_true, Y_pred_classes)
print("Confusion Matrix is :\n"+str(confusion_mtx))
```

Confusion Matrix is :

```
[[886  2 18 13  4  1 70  0  6  0]
```



```
[ 2 983  0  9  3  0  2  0  1  0]
[25  1 880  4 50  0 38  0  2  0]
[11  5 11 901 44  0 26  0  2  0]
[ 3  1 52 18 895  0 30  0  1  0]
[ 0  0  0  0  0 976  0 18  0  6]
[111  1 72 19 68  0 725  0  4  0]
[ 0  0  0  0  0  8  0 973  0 19]
[ 4  1  1  3  4  1  2  4 980  0]
[ 1  0  0  0  0  7  0 27  0 965]]
```

## 13 Incorrect Predictions

```
In [15]: # Errors are difference between predicted labels and true labels
errors = (Y_pred_classes - Y_true != 0)
incorrect = (errors*1).sum()
print ("Number of Incorrect Predicitons are: " + str(incorrect)+ " out of " + str (errors.sum()))
Y_pred_classes_errors = Y_pred_classes[errors]
Y_pred_errors = Y_pred[errors]
Y_true_errors = Y_true[errors]
X_val_errors = testImages[errors]

def display_errors(errors_index,img_errors,pred_errors, obs_errors):
    """ This function shows 6 images with their predicted and real labels"""
    a = ['T-shirt','Trouser','Pullover','Dress','Coat','Sandal','Shirt','Sneaker','Bag']
    n = 0
    nrows = 2
    ncols = 3
    fig, ax = plt.subplots(nrows,ncols,sharex=True,sharey=True)
    for row in range(nrows):
        for col in range(ncols):
            error = errors_index[n]
            ax[row,col].imshow((img_errors[error]).reshape((28,28)))
            ax[row,col].set_title("Predicted label :{}\nTrue label :{}".format(a[pred_errors[n]],a[obs_errors[n]]))
            n += 1
    fig.tight_layout()

# Probabilities of the wrong predicted numbers
Y_pred_errors_prob = np.max(Y_pred_errors,axis = 1)

# Predicted probabilities of the true values in the error set
true_prob_errors = np.diagonal(np.take(Y_pred_errors, Y_true_errors, axis=1))

# Difference between the probability of the predicted label and the true label
```

```

delta_pred_true_errors = Y_pred_errors_prob - true_prob_errors

# Sorted list of the delta prob errors
sorted_dela_errors = np.argsort(delta_pred_true_errors)

# Top 6 errors
most_important_errors = sorted_dela_errors[-6:]

# Show the top 6 errors
print()
print("Displaying the top 6 errors")
display_errors(most_important_errors, X_val_errors, Y_pred_classes_errors, Y_true_errors)

```

Number of Incorrect Predicitons are: 836 out of 10000 inputs

Displaying the top 6 errors

